

CP probe (replay-only, no retrain): desaturating node values via $\tanh(\text{cp_order}/T)$ shifts the scale-invariant headroom metric significantly at the grid extremes -- $T=100$ (MORE saturated than WDL baseline) gives SingleHalt* significantly MORE recovered value (less headroom); $T=600$ (much less saturated) gives significantly LESS recovered value (more headroom) -- motivating CP-RETRAIN.

