


S2: encoder trained to predict tree stats directly [STRUCTURAL-ONLY (the real S2/S3 test)]
targets={n_nodes, height, width} -- new TreeEncoder, MSE regression head, same backbone as E1



Plain-English finding: this curve shows whether an encoder can learn to reconstruct {n_nodes, height, width} when TRAINED for it directly (unlike E1's child-WDL objective, which was shown NOT to).